

Where is the web still insecure?

Regional scans for HTTPS certificates

Anushah Hossain¹, Kristina Nelson¹, and Tjerand Silde^{1,2}

¹ University of California Berkeley, Berkeley, California, USA
{anushah.h, krisn, tjerand.silde}@berkeley.edu

² Norwegian University of Science and Technology, Trondheim, Norway
tjerands@stud.ntnu.no

Abstract

To better understand web security as it is experienced around the world, we scan the top 500 most visited sites of internet users from nine countries of interest. We document HTTPS usage, the encryption algorithm, and certificate information, including issuing date and length of validity. Insights from this project will help benchmark web security prior to upcoming browser interventions[3], and identify regions that may benefit from capacity building to implement TLS in the future.

1 Introduction

HTTPS is an extension of the file transfer protocol HTTP that uses Transport Layer Security (TLS) to ensure that the connection between computers is secured against eavesdroppers and that information sent remains untampered. Sites have been rapidly moving to HTTPS from plain HTTP in the past two years, though most recent internet-wide scans find that approximately 41.5% of active websites are still insecure [6].

In this paper, we highlight two areas worth heeding in ongoing efforts to secure the web. We first consider regional variation in HTTPS usage amongst top sites browsed by international users. Who is mostly likely to encounter sites that are not secure? To answer this, we scan certificates of the top five hundred most visited sites from nine countries of interest. From the certificate data gathered, we are also able to assess the effectiveness of browser policy changes to incentivize implementation of HTTPS, which we discuss in the second half of the paper.

Our findings suggest that parts of the web remain insecure, and that the risk is not experienced evenly across the globe. We do find, however, that browser-led “nudges” may encourage uptake of HTTPS overall.

2 Background

Google and Mozilla both recently announced that they will mark websites explicitly as not secure within the browser address bar if they do not have a valid X.509 certificate. This follows a series of actions taken since 2014 to incentivize HTTPS usage:

1. August 2014: Google made HTTPS-status a ranking signal [4] for internet searches.
2. September 2016: Both Google [9] and Mozilla [7] announced that from January 2017, they will label HTTP pages with password or credit card form fields as “not secure,” given their particularly sensitive nature.
3. February 2018: Google announced that from mid-July 2018, Chrome will mark all HTTP sites as not secure.

A November 2015 report by the National Institute of Standards and Technology (NIST) specifies the set of cryptographic algorithms and key sizes considered “secure” [8]. Acceptable hash-functions are: SHA-256, SHA-384, SHA-512 and SHA-3. RSA is secure with 2048 bit keys and Elliptic Curves (EC) with 224 bit keys. SHA-1 was broken in 2005 and should not be used [1]. RSA with 1024-bit keys is also considered breakable by an adversary with sufficient computational power [5]. We evaluate website security in accordance with these metrics in the following sections.

3 Methods

We select nine countries – Canada, China, Germany, Ghana, India, Iran, Norway, Russia and USA – that range in geography, income level, and political regime. For each country, we scrape the top five hundred most visited sites from the Alexa top sites service. Scans of these lists portray the security of the web as experienced by these countries internet users, but do not necessarily represent the security of locally-hosted sites.

To extract information on website certificates, we use the OpenSSL python library [10]. We see the X.509 certificate if it exists, and record information about the certificate issue and expiration dates, signing algorithm, encryption algorithm, key sizes and HTTP Strict Transport Security (HSTS) usage. HSTS is a policy mechanism where the browser refuses to set up a connection unless HTTPS is used. Determining which sites implement HSTS is made challenging by the variety of ways sites are able to deliver HSTS. Chrome and other major browser vendors have begun shipping a hard-coded preload list of HSTS websites, so that the browser treats them as HSTS without ever having to receive the header. As a result, the sites we identify as using HSTS may only be a fraction of those that actually implement strict transport security.

We collected the site listing data from Alexa on March 26, 2018, and extracted certificate information from the sites on April 14, 2018. We evaluated both the hash function and key length used to better characterize website security.

4 Results

4.1 Regional differences

Figure 1 and 2 summarize our results on site security as experienced by users from each country. We include a Global 500 column of the top sites overall, as ranked by Alexa top sites, as a point of comparison for the country results. We report HSTS percentages as a fraction of total HTTPS traffic, but note that these results are likely to be inaccurate as a result of the different ways by which HSTS is implemented.

Our results show that the United States (87%), Norway (85%), Canada (82%) and Germany (81%) have the highest percentages of top sites using HTTPS. China has the lowest fraction with only 54%. We find that some of the countries with the lowest HTTPS usage have relatively high HSTS usage within our sample. In Iran, for example, only 59% of the most visited countries use HTTPS, but 81% of those sites use HSTS. We hypothesize that this is due to differences in local and foreign websites. If the majority of HTTPS sites visited by internet users from Iran are sites belonging to global companies, for example, and those global sites tend to use HSTS, that would be reflected in a higher relative HSTS rate in our data for Iran. To investigate this in the future, we will need to filter websites by country of origin; we plan to begin testing this hypothesis by comparing overlap across country top 500 listings.

		Global 500	Canada	China	Germany	Ghana	India	Iran	Norway	Russia	USA
HTTPS		85%	82%	54%	81%	63%	67%	59%	85%	64%	87%
HSTS (% of HTTPS sites)		91%	67%	74%	62%	69%	69%	81%	64%	73%	63%
Signing Algorithm	ecdsa-with-SHA256	12%	11%	3%	12%	17%	18%	9%	14%	7%	12%
	sha1WithRSAEncryption	0%	0%	1%	0%	0%	0%	5%	0%	1%	0%
	sha256WithRSAEncryption	88%	89%	95%	88%	83%	81%	86%	85%	92%	88%
Encryption Algorithm, Key Size	EC256	17%	14%	10%	15%	20%	20%	11%	17%	10%	15%
	RSA2048	79%	84%	87%	78%	76%	77%	83%	75%	82%	83%
	RSA4096	3%	2%	2%	6%	4%	2%	4%	7%	7%	2%
Average Certificate Length (months)		19	20	23	20	18	18	36	21	20	19
Total Site Count		500	500	500	500	500	500	500	500	500	500

Figure 1: Results per country.




Figure 2: Security parameters.

We find that RSA with SHA256 is the most common signing algorithm used, followed by ECDSA with SHA256. Only a small fraction of the websites used encryption with insecure algorithms and key sizes. From our combined pool of top sites across countries, twenty-six unique websites still use SHA1 for signings and nine websites still use RSA with key size 1024 bits. Iran has the highest percentage of visited sites using the SHA1 hash function, possibly reflecting insecure local content. Neither RSA with SHA384 nor RSA with SHA512 are used in our sample sites.

Comparison of the Global 500 figures against those of each country show that the experiences of many internet users would be masked by aggregated figures. We see that the majority of countries browse fewer HTTPS sites than the global 500 would suggest, and far fewer HSTS sites.

4.2 Certificate Issuance

Here, we review the certificate data gathered from all of the unique sites scanned (3221 unique sites out of 4500 total scanned). We find that the certificates are typically issued for one, two, or three years, though we do see shorter periods of three or six months. Several certificates issued in China and Iran were valid for 30 to 100 years, but these were marked as insecure in major browsers.

The graphs in Figure 3 visualize the number of certificates that have been issued, how many are expiring soon and how long they are valid after the issue date. We mark the dates when Google and Mozilla released browser security announcements in the left-most figure. As Google and Mozilla begin rolling out strict warnings for insecure websites, we see that an increasing number of websites are issued new certificates.

The sharp increase in new certificates issued after the browser announcements suggests that many websites recently updated their certificates or were issued their first one. Future analysis will account for coinciding factors that may also have contributed to these trends, but our results thus far support the claim that Google and Mozilla’s initiatives are successfully impacting the security culture among the top websites.




Figure 3: Certificate overview.

5 Limitations

Our work is necessarily limited by the sample of sites we examine. As Scheitle et al. [2] document, it is unclear how representative Alexa listings are of the entire web, as they are based on data collected from opt-in browser extensions. Though we considered alternate site listings provided by Quantcast and Majestic Million, we settled upon Alexa listings due to their stratification by country. Other listings typically require paid subscriptions.

We manually verified a subset of websites with valid certificates and websites with errors, and found both false positive and false negative results compared with the Python script. We plan to test whether a purely OpenSSL solution would give different results on the same top 500 lists. It is also possible to use other programming languages such as JavaScript. We hope these next steps can give us both more accurate data and also more insight into which algorithms are used for the symmetric key cryptography component of TLS, and the lengths of their security parameters.

Future work will attempt to distinguish between locally generated content and foreign content to better understand local capacity for implementing TLS, and assess the relationship between site popularity and security. We will also study HSTS more in detail to achieve more accurate results, in addition to check the availability of certificate revocation information and the certificate transparency for the different web pages of interest.

6 Conclusion

We scanned the top five hundred websites accessed by internet users from nine countries to paint a disaggregated portrait of web security. This contrasts with existing scans that evaluate the entire address space or specific sectors. Our results show significant regional variation and suggest that users from China, Ghana, Iran, and Russia are relatively more susceptible to eavesdropping or corrupted data when sending information over the internet. Future work will expand the sample of countries considered to assess broader regional patterns in security. These initial results suggest that while web security is improving, the benefits are not yet evenly distributed globally. Knowledge of where the web is insecure, as experienced by a country's users, can help policy makers and other stakeholders place targeted pressure on the sites in question to implement HTTPS and HSTS, or recommend stronger encryption algorithms.

Thus far, the majority of action incentivizing web security has come from private sector actors, as we've seen in the success of browser policies and cost-decreasing initiatives such as Let's Encrypt. As Google and Mozilla begin to mark websites without HTTPS as insecure, it will be interesting to see how these results evolve. We expect the percentage of websites using HTTPS to increase significantly in the coming months, given past responsiveness to browser policy.

References

- [1] CWI and Google. Shattered. shattered.io, 2017.
- [2] Scheitle et al. Structure and stability of internet top lists. arxiv.org/pdf/1802.02651.pdf, 2018.
- [3] Google. Say yes to https: Chrome secures the web, one site at a time. blog.google/topics/safety-security/say-yes-https-chrome-secures-web-one-site-time, 2011.
- [4] Google. Hhttps as a ranking signal. webmasters.googleblog.com/2014/08/https-as-ranking-signal.html, 2014.
- [5] Matthew Green. How does the nsa break ssl? blog.cryptographyengineering.com/2013/12/03/how-does-nsa-break-ssl/, 2013.
- [6] Qualys Inc. Ssl pulse: Monthly scan. www.ssllabs.com/ssl-pulse, 2018. [Accessed April 2018].
- [7] Mozilla. Communicating the dangers of non-secure http. blog.mozilla.org/security/2017/01/20/communicating-the-dangers-of-non-secure-http, 2017.
- [8] NIST. Transitions: Recommendation for transitioning the use of cryptographic algorithms and key lengths. dx.doi.org/10.6028/NIST.SP.800-131Ar1, 2015.
- [9] The Chromium Projects. Marking http as non-secure. www.chromium.org/Home/chromium-security/marking-http-as-non-secure, 2016.
- [10] The pyOpenSSL developers. pyopenssl documentation. pyopenssl.org/en/stable, 2018.